



AI reinvention made real

Accenture Innovation Challenge 2026

Team details

TEAM NAME: ERROR 404



Saugata Malakar (Team Leader)

College: IIT Kharagpur
Stream: Civil
Year of graduation: 2028



Trinakshi Saha

College: IIT Kharagpur
Stream: Electrical
Year of graduation: 2028



Satadru Sen

College: IIT Kharagpur
Stream: Bioscience and Biotechnology
Year of graduation: 2028

PROBLEM STATEMENT · BusinessIntelligence.ai — The Problem

Dashboards tell you WHAT happened. They rarely tell you WHY — or what to do next.

A regional revenue KPI drops 8%. Today, explaining it is a slow, manual, unscored process that runs independently for every metric a business tracks.

▼ 8%

Revenue

East Region · Week of Aug 11

2–5 DAYS

To diagnose one anomaly

Manual work, one analyst, no checklist

100s

KPIs tracked per business

Too many to review by hand each week

0

Root-cause explanations

Generated by the dashboard on its own

One flagged KPI, one region, one week — and it was already flagged by a routine threshold alert. Multiply this across every metric a business tracks, and the backlog of unexplained shifts never clears.

DOESN'T scale

One analyst, hundreds of KPIs — most anomalies quietly go unreviewed every week.

TOO SLOW

By the time a 2–5 day explanation lands, the business has already moved on.

NO EVIDENCE TRAIL

Dashboards surface correlation, not a tested, provable cause an analyst can trust.

TODAY'S PROCESS · What happens when Revenue drops 8% right now

SOURCES SCANNED

Transactions, marketing spend, support tickets, competitor news — each checked one by one, on no shared timeline

METHOD

Ad-hoc, manual correlation hunting by a single analyst, with no fixed process or checklist to follow

OUTPUT

A best-guess explanation, with no confidence score and no evidence trail attached to it

TIME COST

2–5 days per flagged KPI, before any action can be taken — while the metric keeps drifting

SCALE PROBLEM

Repeated from scratch, quarter after quarter, for every one of the hundreds of KPIs a business tracks

FEEDING THE INVESTIGATION — checked manually, one at a time

Transactions

Marketing

Support Tickets

Competitor News

No shared timeline links them — the analyst has to line them up by hand

WHAT TODAY'S DASHBOARD ACTUALLY GIVES YOU

- ✓

A number & a trend line

the anomaly itself
- ✗

A verified root cause

left to the analyst to find
- ✗

A confidence level

no sense of how sure to be
- ✗

A ranked next action

no guidance on what to do



- ⚠

No evidence trail:

Nothing ties the number to a provable, checkable cause
- 🔄

No root-cause ranking:

Every hypothesis looks equally plausible, unranked
- 🕒


No time saved:

Every KPI restarts the search from zero, every time

The problem isn't a lack of data. It's the lack of an explanation layer.

OUR CONCEPT · The KPI Storytelling Engine — Live in Action

An AI pipeline that turns a metric shift into a grounded explanation and a ranked next step — in minutes, not days. The LLM narrates; it never decides the cause.



Signal: STL • change-point • materiality **Root cause:** Driver tree • causal checks **Narration:** Grounded LLM • evidence validation **Action:** Curated action library

REAL CHANGE, NOT NOISE
Every flagged move clears a significance & materiality test first.

EVIDENCE, NOT CORRELATION
A cause must precede and be proportionate to the shift before ranking.

HONEST WHEN UNCERTAIN
Thin or conflicting evidence is surfaced as-is, or routed to an analyst.

LIVE KPI STORY · East Region Revenue — illustrative example from the problem statement

FROM ALERT → VERIFIED STORY → ACTION

WHAT CHANGED

Revenue ▼ 8% • Purchase Frequency ▼ 9%

EVIDENCE

Checkout-error support tickets ▲ 3×, starting 2 days before the shift

CROSS-CHECK

A comparable, unaffected region shows no similar decline

CONFIDENCE

HIGH — evidence is strong and consistent

RECOMMENDED ACTION

Escalate checkout issue to Engineering; customer-success outreach to affected accounts

TODAY

2–5 DAYS

KPI STORYTELLING ENGINE

→ MINUTES

WHEN EVIDENCE IS AMBIGUOUS, THE ENGINE GRADES ITS OWN CERTAINTY

HIGH → Recommend

MODERATE → Review

LOW → Investigate

Evidence-linked

Every claim traces to a real data point


Every KPI

Not just the few analysts can hand-check

Action, not alert

Ends in a ranked, costed next step

LLM NARRATES. IT NEVER DECIDES THE CAUSE.



Thank you

